PERSON FINDER · A COMPLETE WALKTHROUGH

Finding a person from a single sentence

How a text-and-image embedding model was measured, fine-tuned and turned into a live video tool that boxes the person you describe, explained step by step in plain words.

40 → 71%
first answer correct, before and after fine-tuning

≈ 26 fps
live video, end to end, on one GPU

0.72 vs 0.40
the right shopper against the next-best, in a mall crowd

BEFORE YOU START

How to read this document

Every step answers three questions: **what** was done, **when** in the project it happens, and **why** it was done that way and not some other way. Wherever a technical word first appears, a box called *"In simple words"* explains it without jargon.

> **IN SIMPLE WORDS**
> The project teaches a computer to do what a security guard does when a parent says, "I lost my son, he's wearing a red jacket and black trousers". The computer looks at every person in the photos or the video and points at the one who fits best.

01 The problem
02 The data
03 The model and embeddings
04 Step 1 · Measure before changing anything
05 How the grades work
06 Step 2 · Fine-tuning
07 Step 3 · The results, and how they compare
08 Step 4 · The image-search website
09 Step 5 · Live video
10 Step 6 · A real crowd: the Mall dataset
11 Limits and next steps
12 Glossary
13 Run it yourself

01 · THE PROBLEM

Searching for people with words

Cameras record thousands of people an hour. When someone is lost or needs to be found, nobody has a photo of them to search with, only a description.

This task is called **text-based person retrieval**. The input is a sentence such as *“a woman with long black hair in a white coat and blue jeans”*, and the output is a ranked list of person photos with the best match first. A second version of the same task works on video: the input is a sentence and a live camera feed, and the output is a box around the right person in every frame.

WHY WORDS AND NOT A PHOTO?

Searching with a photo of the person (re-identification) only works when you already have a photo. A witness or a parent only has memory and words.

WHY IS IT HARD?

The same red coat looks dark red in the shade and almost black at night. Descriptions are short and vague, and many people wear similar clothes.

> **IN SIMPLE WORDS · RETRIEVAL**
> “Retrieval” means fetching the right items from a big pile. A search engine retrieves web pages, and this project retrieves people.

02 · THE DATA · *WHEN:* BEFORE ANY CODE

RSTPReid: 4,101 people, 41,010 sentences

The project uses the public **RSTPReid** dataset: 20,505 photos of 4,101 different people taken by 15 street cameras. Each photo comes with two English descriptions written by two different people.

The split, and why it is done by person

SPLIT	PEOPLE	PHOTOS	USED FOR
train	3,701	18,505	teaching the model
val	200	1,000	picking the best saved version during training
test	200	1,000	the final, reported grade. It is touched only once

Why split by person? If the same person appeared in both training and testing, the model could memorise their face and clothes, and the test would look easier than real life. Because nobody in the test set was ever seen during training, the test measures the real skill: finding *strangers*.

> **IN SIMPLE WORDS · TRAIN / VAL / TEST**
> Think of a school year. The **train** set is the homework the student practises on. The **val** set is the mock exam that tells the teacher which week the student was at their best. The **test** set is the final exam, opened once at the end.

03 · THE MODEL · *WHEN:* CHOSEN AT THE START

EmbeddingGemma-2, and what an “embedding” is

The model is Google's **EmbeddingGemma-2**. It does one thing: it turns a piece of text, *or* an image, into a list of 768 numbers. The same model handles both kinds of input.

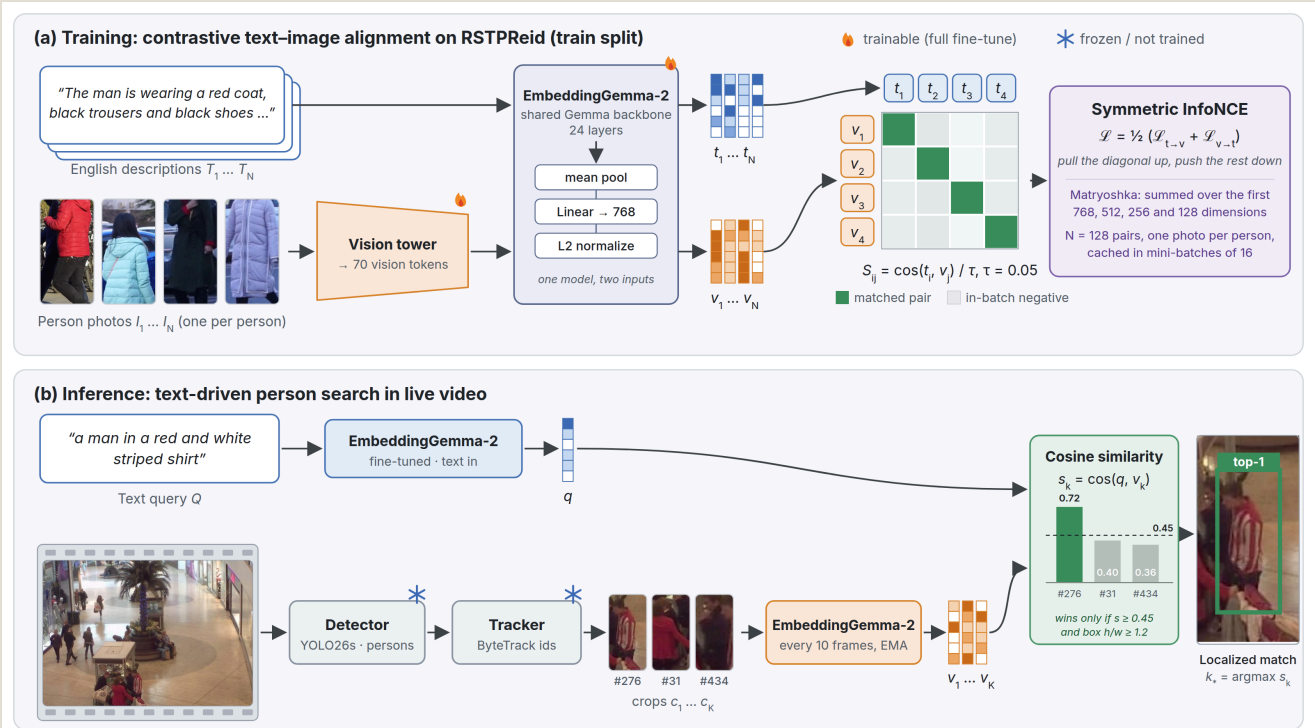
> **IN SIMPLE WORDS · EMBEDDING**
> An embedding is like GPS coordinates for meaning. A sentence about a red jacket and a photo of a red jacket get coordinates that are close together. A sentence about a white coat gets coordinates far away. Finding the right person then becomes “find the nearest point on the map”.

How closeness is measured: cosine similarity

Every list of 768 numbers is scaled to length 1, which is called *L2-normalised*. Two such lists can then be compared by the angle between them. The **cosine** of that angle is 1 when they point the same way, 0 when they are unrelated, and negative when they point apart. Ranking a gallery means computing one cosine per photo and sorting.

Three settings that matter, and why

SETTING	VALUE	WHY
Number format	`bfloat16`	Uses half the memory of normal 32-bit numbers, so the model runs fast on the GPU. The other 16-bit format, float16, is not supported by this model.
Text prompt	`SearchQuery`	The model was trained to treat text differently depending on the job. This prompt tells it “this sentence is a search”, which ranks better than plain text.
Vision tokens	`70`	An image is cut into small pieces called “tokens”. The model allows 70 to 1,120 tokens. Person crops are small and blurry, so 70 was both the fastest and the most accurate in the sweep (280 was the worst).



(a) Training: descriptions and photos pass through the same EmbeddingGemma-2 backbone, and the loss pulls each matching pair together. (b) Inference in video: detected people are cropped, encoded and compared with the query, and the best match is boxed.

04 · STEP 1 · *WHEN:* FIRST, BEFORE ANY TRAINING · `TBPR_EX.IPYNB`

Measure the model as it is: the zero-shot baseline

You cannot know if training helped unless you know where you started. So the first experiment uses the model exactly as Google released it.

> **IN SIMPLE WORDS · ZERO-SHOT**
> "Zero-shot" means the model is tested on a task without any practice on it, like asking someone to referee a sport they have only read about.

1. **Encode all 20,505 photos once** and save them to `outputs/image_embeddings.npy`. This takes about 15 minutes. *Why save them?* Photos don't change, so there is no reason to recompute them for every search. The website reuses this file.
2. **Encode all 41,010 descriptions** with the `SearchQuery` prompt.
3. **Rank**: for every description, compute the cosine with every photo on the GPU, in chunks so that memory never overflows, and record where the right person landed.

ZERO-SHOT, TEST SPLIT	RANK-1	RANK-10	MAP
280 vision tokens	35.5	70.0	25.0
70 vision tokens	**40.1**	**75.4**	**28.8**

What this taught us. Out of the box, the model already finds the right person first 40% of the time among 1,000 photos, which is a real but modest skill. Searching the *whole* dataset (20,505 photos) drops Rank-1 to 10%. Re-running with random galleries of growing size showed why: Rank-1 falls steadily from 46% at 250 photos to 10% at 20,505. The task gets harder with more people to confuse, not because some photos are worse.

> **WHY THIS ORDER**
> Measuring first fixes a fair reference point, with the same split, the same metrics and the same code. Every later number is compared to it.

05 · THE GRADES

Rank-1, Rank-5, mAP, mINP, without the maths

GRADE	QUESTION IT ANSWERS	PLAIN EXAMPLE
Rank-1	Is the first photo the right person?	The guard points once. Right or wrong?
Rank-5 / Rank-10	Is the right person among the first 5 / 10?	The guard shows a short list, and the parent spots their child.
mAP	Are *all* photos of that person near the top?	Not just one sighting: every camera that saw them.
mINP	How deep is the *hardest* photo of that person?	The one blurry, dark sighting. Is it still found?

The test runs each of the 2,000 test descriptions against the 1,000 test photos (**protocol A**). Every grade is an average over those 2,000 searches. All grading code lives in one shared file, `tbpr.py`, so the zero-shot and fine-tuned notebooks are graded in exactly the same way.

> **WHY SEVERAL GRADES?**
> One number hides things. A model can be good at the first guess (Rank-1) but leave the person's other photos buried (low mAP). Reporting all of them shows both.

06 · STEP 2 · *WHEN:* AFTER THE BASELINE · `TBPR_FINETUNE.IPYNB`

Fine-tuning: teach the model what matters for people

The original model learned from general images and text on the internet. It has never been told that “dark red hooded top” and a blurry street crop of a red coat are the same person. Fine-tuning shows it 37,010 examples of exactly that.

> **IN SIMPLE WORDS · FINE-TUNING**
> Fine-tuning is extra practice on one specific job. The model keeps everything it already knows, and its numbers are nudged a little at a time until it gets better at this job.

The loss: a multiple-choice test, taken both ways

Each step takes 128 (sentence, photo) pairs. For each sentence, the correct photo is the one it was written for, and the other 127 photos are wrong answers. The model is rewarded for giving the right pair the highest score in its row. The same test is then run the other way: each photo must pick its sentence. This is **contrastive learning**, using `CachedMultipleNegativesRankingLoss`.

- **Why 128 pairs?** More wrong answers per question make a harder, more useful test.
- **Why “cached”?** 128 images at once do not fit in 24 GB of GPU memory. The cached version computes them 16 at a time and combines the results, which gives the same maths with less memory.
- **Why both directions?** Search goes text to image, but training image to text as well makes the shared space more balanced.

The batch rule: never two photos of the same person

Each person has about five photos. In a random batch of 128, the same person sometimes appears twice, and the loss would then call a *correct* match “wrong”. A small custom sampler, `UniquePersonBatchSampler`, builds every batch from 128 *different* people. **Why it matters:** without it the model would be punished for being right, which slowly teaches it the wrong thing.

Matryoshka: one vector, four sizes

The loss is also applied to the first 512, 256 and 128 numbers of each vector on their own (`MatryoshkaLoss`). The model learns to put the most useful information first, like Russian nesting dolls. **Why:** a 256-number index is 3× smaller and faster, and after fine-tuning it loses only about one Rank-1 point.

Settings, and the reason for each

SETTING	VALUE	WHY
learning rate	2e-5	Small steps, so the model improves at this task without forgetting what it already knows.
warm-up	first 10% of steps	Starts even gentler, because the first updates are the noisiest.
epochs	3 (≈ 33 min each)	The validation grade stops improving after epoch 1, so more epochs would only waste time.
model selection	best val mAP	The test set is never used to choose, so the final grade stays honest.
GPU memory	≈ 18 GB	It fits on one RTX 3090 (24 GB).

07 · STEP 3 · *WHEN:* AFTER TRAINING, ONCE

The results

Fine-tuning nearly doubles every grade. **678** test searches moved up to first place, and 51 dropped out of it. The right person is now in the top 10 for 93% of searches.

Why it works: measured on six unseen test people, before training a sentence was about equally close to everyone (0.73 to its own person, 0.66 to others). After training, the gap is wide (0.51 against 0.12). That wide gap is what lets a single threshold separate the right person from the crowd in video later.

07B · COMPARISON · *WHEN:* LAST, AFTER EVERY EXPERIMENT · TBPR_COMPARE.IPYNB

How it compares with CLIP, IRRA, RDE and others

A number like 71% only means something next to other numbers. So the model was compared two ways: with famous general models run on exactly the same test, and with the published results of methods built for this task.

> **IN SIMPLE WORDS · CLIP AND SIGLIP**
> CLIP (OpenAI) and SigLIP (Google) learned from hundreds of millions of internet photos with their captions. They are the standard "text meets image" models. Here they are used **zero-shot**: as downloaded, with no practice on person search.

> **IN SIMPLE WORDS · IRRA, RDE, RASA, APTM**
> These are research methods built for exactly this task and trained on the same RSTPReid training people. Their numbers are copied from their papers, because they were not re-run here. Most of them start from a CLIP-sized image model (ViT-B/16), which is smaller than EmbeddingGemma-2.

On RSTPReid (photos and long descriptions)

METHOD	TRAINED ON RSTPREID?	RANK-1	RANK-10	MAP	MINP
EmbeddingGemma-2, fine-tuned (this project)	yes	**71.1**	**93.2**	**54.7**	**29.5**
APTM (ACM MM 2023)	yes	67.5	91.5	52.6	–
RaSa (IJCAI 2023)	yes	66.9	91.4	52.3	–
RDE (CVPR 2024)	yes	65.4	89.9	50.9	28.1
IRRA (CVPR 2023)	yes	60.2	88.2	47.2	25.3
CLIP ViT-B/16, fine-tuned (AAAI 2024 baseline)	yes	56.2	86.6	43.3	–
EmbeddingGemma-2, zero-shot	no	40.0	75.3	28.8	9.7
SigLIP B/16, zero-shot	no	30.5	67.0	22.0	6.5
CLIP ViT-L/14, zero-shot	no	11.3	37.7	8.2	2.0
CLIP ViT-B/16, zero-shot	no	9.9	34.4	7.2	1.6

What this shows. After fine-tuning, the model puts the right person first more often than every listed method, and among the methods that report mINP, it finds even the hardest photo of a person best. Zero-shot CLIP manages only about 10%. RSTPReid descriptions are long, and many people wear similar dark clothes, so a model has to read details such as "glasses" or "a pink bag strap", and general models were never taught to do that for people. Rank-1 here is 71.1. The training notebook measured 71.4, because re-encoding with 16-bit numbers moves results by about 0.3.

On live video (the Mall crowd, short colour descriptions)

How it was tested. The detector and tracker ran once, and every model scored exactly the same people. For four descriptions, everyone who fits was marked by eye. The test then asks two questions: when such a person is in view, is the best-scoring person a right one (**top-1**)? And how cleanly can one threshold separate right from wrong people (**AUC**, where 1.0 is perfect)?

EMBEDDER	TOP-1 IN VIEW	AUC	STRIPED SHIRT	CROPS / SECOND
EmbeddingGemma-2, fine-tuned	**86.2%**	**0.952**	**96.8%**	34
EmbeddingGemma-2, zero-shot	85.2%	0.935	93.5%	34
CLIP ViT-B/16, zero-shot	84.3%	0.949	86.4%	345
CLIP ViT-L/14, zero-shot	82.4%	0.936	87.7%	143
SigLIP B/16, zero-shot	74.7%	0.916	64.3%	345

What this shows. With short colour cues such as “a woman in a red coat”, the gap almost closes. CLIP has seen millions of red coats, and colour alone often picks the right person. The fine-tuned model is still the best overall, and clearly the best on the patterned “red and white striped shirt”, but CLIP encodes crops 10× faster. **Why it matters:** the video demo stays real-time because each person is re-encoded only every 10 frames, and that is what makes the slower but more accurate model affordable. This test is small (4 descriptions, 573 frames), so read it as a trend, not a ranking.

08 · STEP 4 · *WHEN:* AFTER THE MODEL IS FINAL · `DEMO/SERVER.PY`

The image-search website

A small web server, written with Python's built-in HTTP library and no web framework, loads the fine-tuned model and the pre-computed photo embeddings. When you type a sentence, it encodes only that sentence (a few milliseconds), scores it against all photos with one matrix multiplication on the GPU, and returns the top 5.

- **Why pre-compute the photos?** Encoding 20,505 photos takes minutes, while encoding one sentence takes milliseconds. Doing the slow part once makes every search instant.
- **Why must the file match the model?** Embeddings from two different models live on two different “maps”. Mixing them would compare GPS coordinates from two different planets.
- **Why only known image names are served:** the server never reads a file path the browser invents, so nobody can ask it for files outside the dataset.

```
GET  /api/example          a random test description and its person id
POST /api/search           {query, gallery: "test" | "all"}  →  top-5 photos + scores
GET  /imgs/<name>          only names that exist in the dataset
```

09 · STEP 5 · *WHEN:* AFTER SEARCH WORKED · `DEMO/TRACK_PERSON.PY`

Live video: from a frame to one green box

The model only understands photos of single people. Video is turned into exactly that: a stream of single-person photos.

STAGE	WHAT	WHY THIS CHOICE
Detect	YOLO26s, “person” class only	Small and fast, and it finds where people are so that the language model only looks at people.
Track	ByteTrack gives each person an id	A person keeps their id across frames, so evidence collected over time is not thrown away.
Re-encode	each track every 10 frames, in one GPU batch	The person rarely changes in a third of a second. This saves about 90% of the work.
Smooth	EMA: `e ← 0.7·e + 0.3·new`	One blurry crop cannot flip the decision, because the track keeps a running memory.
Gate 1	score ≥ 0.45	The right person scored 0.45–0.76, while strangers peaked at 0.40–0.42 when that person was absent. Below 0.45, no one gets the box.
Gate 2	box height ÷ width ≥ 1.2	Training photos are full bodies. Head-only boxes and boxes covering two people score unreliably, so they can never win.

> **IN SIMPLE WORDS · EMA**
> An exponential moving average is a memory with fading weight. Every new look counts for 30%, and everything seen before counts for 70%. Recent looks matter most, but no single look can take over.

Speed. About 26 frames per second end to end on 768×576 video, and 55–100 fps once every track is encoded. In the website, a background thread plays the video in real time and skips frames if it falls behind, so the video never slows down. A new sentence takes effect on the very next frame, about 150 ms after the click.

A real crowd: the Mall dataset

Everything so far was tested on clean, cut-out photos. A shopping mall is the real thing: 13–53 people in each frame, people hiding behind each other, and tiny figures far from the camera.

The Mall dataset (Chen et al., BMVC 2012) has 2,000 frames from one camera. A human put a dot on **every head**, 62,315 dots in total. There are no names or ids, so it can grade one question exactly: *does the detector see the people?* The frames were also turned into a video, `videos/mall.mp4`, which plays in the website.

How the detector is graded

Each head dot may be matched to at most one person box, and only if the dot sits in the top 45% of that box. The optimal one-to-one pairing is found with the Hungarian algorithm. Matched dots count as hits, unmatched dots are misses, and unmatched boxes are false alarms.

DETECTOR SETTING	PRECISION	RECALL	COUNT ERROR / FRAME
640 px, confidence 0.40 (the demo's default)	97.7%	39.0%	18.7
1280 px, confidence 0.25	92.4%	**66.3%**	**8.8**

> **IN SIMPLE WORDS · PRECISION AND RECALL**
> **Precision** asks: when the detector says “person”, how often is it right? **Recall** asks: of all the real people, how many did it find? The demo's detector is almost never wrong, but in a crowd it misses many far-away people.

What this taught us. The language model transfers to a place it never saw: each description finds its shopper among 20–30 people, with best scores of 0.60–0.72. The weak link in a crowd is the detector, because a person who never gets a box can never be chosen. Doubling the detector's input size to 1280 px finds 70% more people.

What it cannot do yet, and what would come next

- **Small and hidden people.** In dense crowds, about a third of people get no box even at 1280 px. The next step would be a detector trained on crowds, or tiling the frame into overlapping pieces.
- **Similar clothes.** Two people dressed in red end up close together in the embedding space. Details such as “glasses” or “scarf” help, but a vague sentence stays ambiguous.
- **Low frame rates.** The Mall camera records under 2 frames per second, so ByteTrack often breaks tracks and the EMA memory restarts. Normal 25–30 fps cameras do not have this problem.
- **Language.** Training used English descriptions only.
- **Responsible use.** This is a research demo on public research datasets. Searching for real people in real footage needs consent, legal grounds and human review of every match.

Every term, in one line

Embedding	A list of numbers (here 768) that captures what a text or image means, so similar things get similar lists.
Cosine similarity	How much two lists point the same way: 1 is the same, 0 is unrelated.
L2-normalised	Scaled to length 1, so only the direction counts.
Zero-shot	Used on a task with no practice on it.
Fine-tuning	Extra training on one specific task, starting from a trained model.
Contrastive loss	Training by “match the right pair, push away the wrong ones”.
Batch	The group of examples the model studies in one step (here 128).
Negative	A wrong answer in the batch that the model should score low.
Epoch	One full pass over the training data.
Learning rate	How big each training step is.
Matryoshka	Training so the first part of a vector works on its own.
Vision tokens	The small pieces an image is cut into before the model reads it.
bfloat16	A 16-bit number format: half the memory, nearly the same accuracy.
Rank-k	Share of searches with the right person in the top k.
mAP / mINP	How early *all* / the *hardest* of a person's photos are found.
Detector (YOLO)	Finds where people are in a frame and draws boxes.
Tracker (ByteTrack)	Gives each box a stable id from frame to frame.
EMA	A running average that favours recent values.
Precision / recall	How often a detection is right / how many real people are found.
CLIP / SigLIP	General text-and-image models trained on internet photos and captions, used here without extra training.
AUC	The chance that a right person scores higher than a wrong one: 1.0 is perfect, 0.5 is a coin flip.

Hungarian algorithm	Finds the best one-to-one pairing between two sets, here boxes and head dots.

Commands

```
# 1. measure zero-shot, fine-tune, stress-test on the mall, then compare (each notebook saves
its outputs)
.venv/bin/jupyter execute --inplace TBPR_ex.ipynb
.venv/bin/jupyter execute --inplace TBPR_finetune.ipynb     # ≈ 18 GB GPU, ≈ 1.7 h
.venv/bin/jupyter execute --inplace TBPR_mall.ipynb
.venv/bin/jupyter execute --inplace TBPR_compare.ipynb

# 2. the website: image search on /, live video on /video
.venv/bin/python demo/server.py --port 8000 \
    --model models/embeddinggemma2-rstpreid --embeddings
outputs/finetuned/image_embeddings.npy \
    --video videos/*.mp4 --video-query "a man in a red and white striped shirt"

# 3. video in a desktop window (type a new sentence in the terminal + Enter)
.venv/bin/python demo/track_person.py --model models/embeddinggemma2-rstpreid \
    --source videos/mall.mp4 --query "a woman in a red coat"

# 4. rebuild the figures in docs/diagrams and this PDF
.venv/bin/python docs/diagrams/build_figures.py
.venv/bin/python docs/diagrams/render.py
.venv/bin/python docs/explainer/build_pdf.py
```